

# A LAZY CONCURRENT LIST-BASED SET ALGORITHM

30 MARCH 2023

**Authors:** Steve Heller, Maurice Herlihy, Victor Luchangco, Mark Moir, William N Scherer III, and Nir Shavit

**Presenter:** Jun-Qing Lim



UNIVERSITY OF  
**WATERLOO**

DAVID R. CHERITON SCHOOL  
OF COMPUTER SCIENCE

# Problem

- List-based sets: basic building blocks for concurrent algorithms
- Current lock-free algorithms (mostly) require the operations with list traversals:
  - To use a mark bit as reference of nodes being logically removed
  - Then perform cleanup – which is expensive (due to overheads on “helping to cleanup”)



# Ways to Synchronize Lists

## 1. Coarse-grained

- single lock
- simple but no concurrency

## 2. “hand-over-hand”

- lock for each successive node before releasing predecessor’s lock
- higher concurrency than (1) but acquire many locks
- in high contention situations, thread 1 waits for thread 2, thread 2 waits for thread 3...
- which slows down the application (expensive)



# New Idea Proposed - “lazy” synchronization

- Based on “optimistic” locking scheme
- 3 main operations:
  - `add(x)`
    - adds x into the set, returns true if x is not already in the set
  - `remove(x)`
    - removes x from the set, return true if x was in the set
  - `contains(x)`
    - returns true if the set contains x




# Algorithm - Overview

- Each operation searches without acquiring locks / interfering others
- Only locks when it locates the position to add/remove → synchronization method!
- In removing, nodes are marked as removed (logical) → then physically unlinked
- Other threads do not “help” in unlinking while traversing (hence “lazy”!)
- Avoids unnecessary overheads

# Algorithm - Design

- Maintain list in increasing order
- Locks when position is located
- Traverse to search for 2.... (for update ops)


```
struct Entry {  
    int key; // value of element  
    Entry next; // reference to the next object  
    bool marked; // logically removed?  
    mutex lock; // used for synchronization  
}
```



# Algorithm - Design


- Maintain list in increasing order
- Locks when position is located
- Traverse to search for 2.... (for update ops)

```
struct Entry {  
    int key; // value of element  
    Entry next; // reference to the next object  
    bool marked; // logically removed?  
    mutex lock; // used for synchronization  
}
```



# Algorithm - add()

1. Traverse to find entry of `curr.key >= key`
2. Locks the predecessor and current
3. Validate to see if the values are changed
4. Perform insert OR restart
5. Always unlocks the entries



```
bool add(int key) {  
    while (true) {  
        Entry pred = head;  
        Entry curr = head.next;  
        while (curr.key < key) {  
            pred = curr; curr = curr.next;  
        }  
        pred.lock();  
        try {  
            curr.lock();  
            try {  
                if (validate(pred, curr)) {  
                    if (curr.key == key) {  
                        return false;  
                    } else {  
                        Entry entry = Entry(key);  
                        entry.next = curr;  
                        pred.next = entry;  
                        return true;  
                    }  
                }  
            }  
        } finally {  
            curr.unlock();  
        }  
    } finally {  
        pred.unlock();  
    }  
}
```

Find position

Locks

Insert



# Algorithm - add()

- Why validate?
  - Gap between unsynchronized traversal and locks
  - pred/curr could be removed
  - Some entry inserted between pred and curr

```
bool validate(Entry pred, Entry curr) {  
    return !pred.marked && !curr.marked && pred.next == curr;  
}
```

```
bool add(int key) {  
    while (true) {  
        Entry pred = head;  
        Entry curr = head.next;  
        while (curr.key < key) {  
            pred = curr; curr = curr.next;  
        }  
        pred.lock();  
        try {  
            curr.lock();  
            try {  
                if (validate(pred, curr)) {  
                    if (curr.key == key) {  
                        return false;  
                    } else {  
                        Entry entry = Entry(key);  
                        entry.next = curr;  
                        pred.next = entry;  
                        return true;  
                    }  
                }  
            }  
        } finally {  
            curr.unlock();  
        }  
    } finally {  
        pred.unlock();  
    }  
}
```


Find position

Locks

Insert

# Algorithm - remove()

1. Traverse to find entry of `curr.key >= key`
2. Locks the predecessor and current
3. Validate to see if the values are changed
4. Perform remove OR restart
  - Logical removal
  - Physical removal
5. Always unlocks the entries



```
bool remove(int key) {  
    while (true) {  
        Entry pred = head;  
        Entry curr = head.next;  
        while (curr.key < key) {  
            pred = curr; curr = curr.next;  
        }  
        pred.lock();  
        try {  
            curr.lock();  
            try {  
                if (validate(pred, curr)) {  
                    if (curr.key != key) {  
                        return false;  
                    } else {  
                        curr.marked = true;  
                        pred.next = curr.next;  
                        return true;  
                    }  
                }  
            } finally {  
                curr.unlock();  
            }  
        } finally {  
            pred.unlock();  
        }  
    }  
}
```

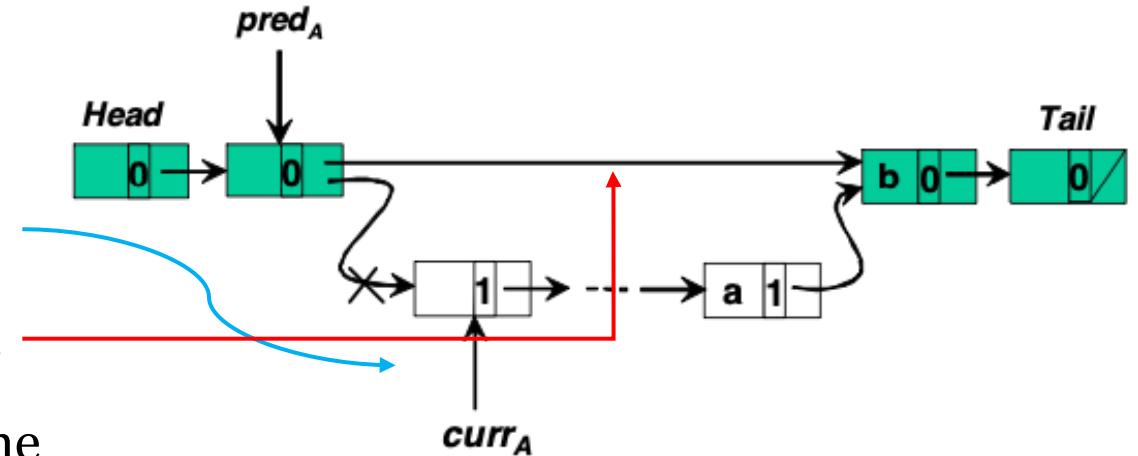
Find position

Locks

Removes

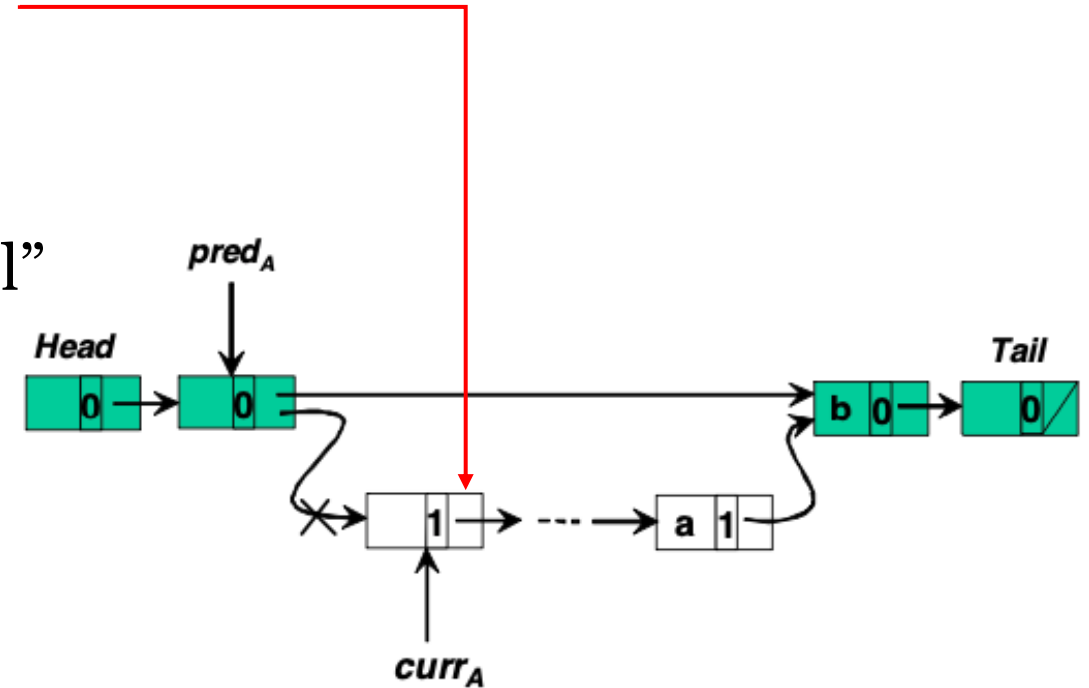
# What happens to concurrent threads traversal?

- Can traverse logically/physically removed nodes
  - Suppose thread A traverses with `pred` and `curr`
  - Then another thread B removes the node `currA`
  - `currA` can still move to the next node because the pointer is still there



# What happens to concurrent threads traversal?


- Say if some thread stops its search at the removed node, and performs some operation
- Operation can still return correct result
- Because nodes are marked as “logical removal”
- Validate checks for !marked



# Algorithm – contains()

- Since node is marked for logical removal
- Also used to determine if it's still present
- Usual traversal + return true if == key and unmarked
- Correct because removed nodes must be marked or not present at all

```
bool contains(int key) {  
    Entry curr = head;  
    while (curr.key < key) {  
        curr = curr.next;  
    }  
    return curr.key == key && ! curr.marked;  
}
```



# Advantage?

1. `contains ( )` is wait-free - does not interfere with any concurrent operations
2. Traversals are not delayed by physical removals in `remove()`
  - Eventually a thread will be able to decide if the node really exists through the marking
- Compared to Michael's lock-free list algorithm
  - When it encounters a logically removed node, the thread “helps in physical removal” with CAS
  - Traversals to be restarted and drop in performance

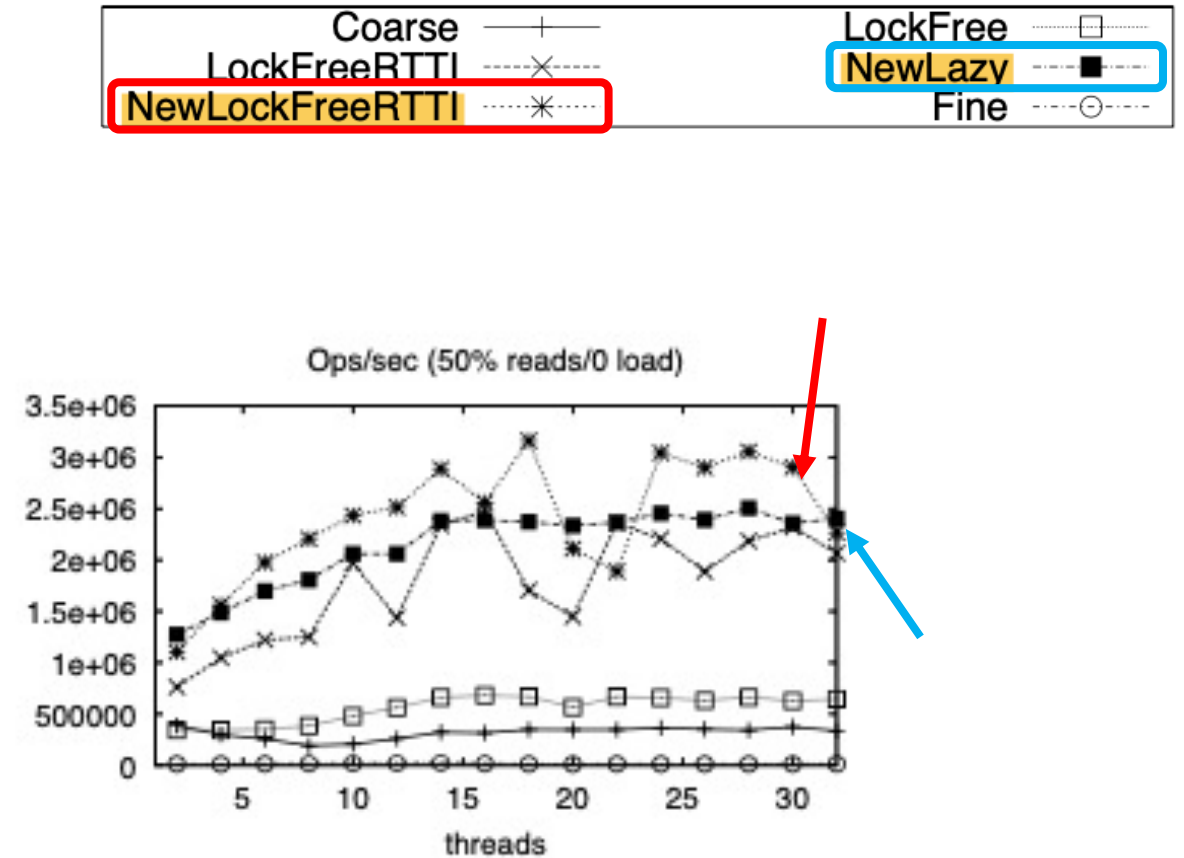
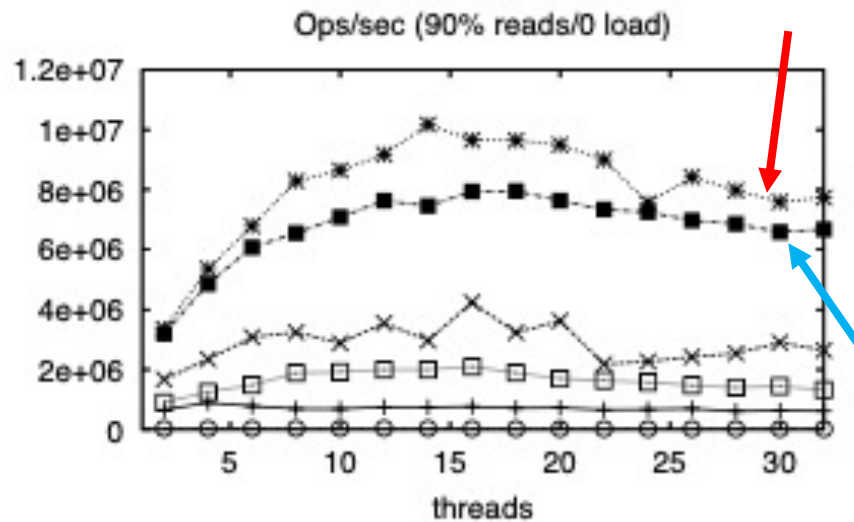
# Performance

- Compared 6 different algorithms with different locking techniques
  1. **Coarse**: single lock
  2. **Fine**: “hand-over-hand” locking (1 lock per item)
  3. **LockFree**: lock-free list using marking but not wait-free for contains() [Michael’s]
  4. **LockFreeRTTI**: improved version of LockFree + Java’s RTTI mechanism

**RTTI**: Run-Time Type Identification – exposes type info during runtime
  5. **NewLockFreeRTTI**: algo (4) but wait-free using the new contains ( )
  6. **NewLazy**: the new algorithm!

# Results?

- Two tests:
  - 90% contains with 10% updates
  - 50% contains with 50% updates
- Both tests outperform by almost double!





# Conclusion

- Paper introduced lazy list – based on lazy marking and deletion of nodes
- Main improvement performance comes from the wait-free searching

**THANK YOU!**

